

# Homework Assignment 3

Parker Reedy

November 26, 2023

1.

```
set.seed(123)

dat <- model.matrix(Sales~., Carseats)[, -1]
train = sample(nrow(dat), 30)
x.train = dat[train, ]
y.train = Carseats[train, ]$Sales

# The rest as test data
x.test = dat[-train, ]
y.test = Carseats[-train, ]$Sales
```


a. The best lambda is 0.02848

```
set.seed(123)

lambda.list.ridge = 1000 * exp(seq(0, log(1e-5), length = 100))

cv.ridge.out = cv.glmnet(x.train, y.train, alpha=0, lambda = lambda.list.ridge)

plot(cv.ridge.out) + abline(v = log(cv.ridge.out$lambda.min), col="red", lwd=3, lty=2)
```



```
## integer(0)
```

```
bestlam = cv.ridge.out$lambda.min
bestlam
```

```
## [1] 0.02848
```

```
out = glmnet(x.train,y.train,alpha=0)
predict(out, type='coefficients', s=bestlam)[1:12,]
```

```
##      (Intercept)      CompPrice      Income      Advertising      Population
##      6.5730326      0.0938637      0.0075634      0.0701387      -0.0001237
##      Price      ShelfLocGood ShelfLocMedium      Age      Education
##      -0.0845304      4.0438779      1.6198782      -0.0527367      -0.0939303
##      UrbanYes      USYes
##      0.5532566      -0.0179076
```

- b. The test MSE is 1.46 and the training MSE is 0.507. This is a fairly low MSE for both training and test. It is possible that we have over fit the data on the training set because the test MSE is 3 times the training MSE. but nonetheless, they are quite good values.

```
ridge.mod = glmnet(x.train,y.train,alpha=0,lambda=bestlam)

ridge.pred=predict(ridge.mod,s=bestlam ,newx=x.test)
mean((ridge.pred-y.test)^2)
```

```
## [1] 1.46
```

```
ridge.pred.train=predict(ridge.mod,s=bestlam ,newx=x.train)
mean((ridge.pred.train-y.train)^2)
```

```
## [1] 0.5074
```

- c. The optimal lambda value is 0.0366. The lasso coefficient estimates corresponding to this lambda value are shown below. There are two coefficients set to zero which are Population, and USYes. This means that in the lasso subset selection, Population and USYes do not have an affect on the response variable Sales because they are effectively excluded from the model.

```
set.seed(123)

lambda.list.lasso = 2 * exp(seq(0, log(1e-4), length = 100))

cv.lasso.out = cv.glmnet(x.train, y.train, alpha=1, lambda = lambda.list.lasso)

bestlamlas = cv.lasso.out$lambda.min
bestlamlas
```

```
## [1] 0.03661
```

```
lasso.mod = glmnet(x.train, y.train, alpha=1, lambda=bestlamlas)

lasso.pred = predict(lasso.mod, s=bestlam, newx = x.train)

lasso.coef = predict(lasso.mod, type='coefficients', s=bestlamlas)
lasso.coef
```

```
## 12 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept)  6.037984
## CompPrice    0.107292
## Income       0.006439
## Advertising  0.075699
## Population   .
## Price        -0.094085
## ShelveLocGood 4.112830
## ShelveLocMedium 1.662509
## Age          -0.057810
## Education    -0.082580
## UrbanYes     0.562462
## USYes        .
```

- d. The training MSE for the lasso model is 0.546 while the Test MSE is 1.464. Similar to the ridge model, it could be the case that we are overfitting our data

```
lasso.pred=predict(lasso.mod,s=bestlamlas ,newx=x.test)
mean((lasso.pred-y.test)^2)
```

```
## [1] 1.464
```

```
lasso.pred.train=predict(lasso.mod,s=bestlamlas ,newx=x.train)
mean((lasso.pred.train-y.train)^2)
```

```
## [1] 0.5466
```

- e. In this application, the Ridge and Lasso methods for estimates are both very similar. They have basically the same test MSE and the training MSE for lasso is only 0.04 more than the ridge training MSE. For this data, It might be better to use the Lasso because it makes the model more sparse which in turn makes it simpler.

2.

```
drug <- read_csv('drug.csv',
col_names=c('ID','Age','Gender','Education','Country',
'Ethnicity','Nscore',
'Escore','Oscore','Ascore','Cscore',
'Impulsive','SS','Alcohol','Amphet','Amyl','Benzos',
'Caff','Cannabis','Choc','Coke','Crack','Ecstasy',
'Heroin','Ketamine','Legalh','LSD','Meth',
'Mushrooms','Nicotine','Semer','VSA'))
```

```
## Rows: 1885 Columns: 32
## -- Column specification -----
## Delimiter: ","
## chr (19): Alcohol, Amphet, Amyl, Benzos, Caff, Cannabis, Choc, Coke, Crack, ...
## dbl (13): ID, Age, Gender, Education, Country, Ethnicity, Nscore, Escore, Os...
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

a.

```
drug$recent_nicotine_use <- factor(ifelse(drug$Nicotine >='CL3','Yes', 'No'), levels = c('No', 'Yes'))
head(drug$recent_nicotine_use)
```

```
## [1] No Yes No No No Yes
## Levels: No Yes
```

b.

```
sub_drug <- drug[, c(2:13, 33)]
colnames(sub_drug)
```

```
## [1] "Age"          "Gender"        "Education"
## [4] "Country"      "Ethnicity"     "Nscore"
## [7] "Escore"       "Oscore"        "Ascore"
## [10] "Cscore"       "Impulsive"     "SS"
## [13] "recent_nicotine_use"
```

c.

```
set.seed(123)

trainsamp = sample(nrow(sub_drug), 1000)

drug.train = sub_drug[trainsamp, ]

drug.test = sub_drug[-trainsamp, ]

head(drug.train)
```

```
## # A tibble: 6 x 13
##   Age Gender Education Country Ethnicity Nscore Escore Oscore Ascore
##   <dbl> <dbl>    <dbl>    <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 -0.952 -0.482   -1.23    0.961   -0.317 -0.580 -1.09  -0.583 -2.08
## 2 -0.952  0.482    0.455    0.961   -0.317  0.313 -0.806  1.06   0.761
## 3  1.09   0.482    0.455    0.961   -0.317  0.826 -0.948 -0.452  0.761
## 4  1.82   -0.482   -1.23    0.961   -0.317 -1.87   0.476 -1.12  -0.0173
## 5 -0.0785  0.482    0.455    0.249    0.126 -0.921  1.29  -0.0193 -0.0173
## 6 -0.952 -0.482   -0.611  -0.285   -0.317  0.417 -0.440  1.06  -1.48
## # i 4 more variables: Cscore <dbl>, Impulsive <dbl>, SS <dbl>,
## #   recent_nicotine_use <fct>
```

d.

```
mod1 <- glm(recent_nicotine_use ~., data=drug.train, family = 'binomial')
summary(mod1)
```

```
##
## Call:
## glm(formula = recent_nicotine_use ~ ., family = "binomial", data = drug.train)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.49460    0.15657   3.16   0.0016 **
## Age         -0.44225    0.09166  -4.82  1.4e-06 ***
## Gender       -0.36206    0.16300  -2.22   0.0263 *
## Education    -0.24006    0.08186  -2.93   0.0034 **
## Country      -0.49745    0.12347  -4.03  5.6e-05 ***
## Ethnicity    -0.22305    0.42973  -0.52   0.6037
## Nscore       -0.11492    0.09117  -1.26   0.2075
```

```
## Escore      -0.10898    0.09443   -1.15    0.2485
## Oscore       0.16778    0.08888    1.89    0.0591 .
## Ascore       0.00979    0.08071    0.12    0.9034
## Cscore      -0.22607    0.08642   -2.62    0.0089 **
## Impulsive    0.10033    0.10357    0.97    0.3327
## SS          0.31365    0.11201    2.80    0.0051 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1366.1 on 999 degrees of freedom
## Residual deviance: 1131.0 on 987 degrees of freedom
## AIC: 1157
##
## Number of Fisher Scoring iterations: 4
```

e.

```
tree.drug <- tree(recent_nicotine_use~., data = drug.train)
tree.drug
```

```
## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
## 1) root 1000 1000 Yes ( 0.43 0.57 )
##    2) SS < -0.06794 455 600 No ( 0.62 0.38 )
##      4) Education < 0.197735 207 300 Yes ( 0.48 0.52 )
##        8) Country < 0.605025 74 90 Yes ( 0.30 0.70 ) *
##        9) Country > 0.605025 133 200 No ( 0.58 0.42 )
##          18) Cscore < -0.20942 51 70 Yes ( 0.35 0.65 ) *
##          19) Cscore > -0.20942 82 100 No ( 0.72 0.28 )
##            38) Age < 0.796185 45 60 No ( 0.56 0.44 ) *
##            39) Age > 0.796185 37 20 No ( 0.92 0.08 ) *
##      5) Education > 0.197735 248 300 No ( 0.73 0.27 ) *
##    3) SS > -0.06794 545 600 Yes ( 0.27 0.73 )
##      6) Country < 0.230255 293 300 Yes ( 0.17 0.83 ) *
##      7) Country > 0.230255 252 300 Yes ( 0.39 0.61 )
##        14) Cscore < 0.671595 197 200 Yes ( 0.32 0.68 ) *
##        15) Cscore > 0.671595 55 70 No ( 0.62 0.38 ) *
```

f. The best size for our tree using 5-fold cross validation is 5

```
set.seed(2)
cv = cv.tree(tree.drug, FUN=prune.misclass, K=5)
cv
```

```
## $size
## [1] 8 7 5 2 1
##
## $dev
## [1] 332 332 327 338 429
```

```
##
## $k
## [1] -Inf 0.0 6.5 15.0 107.0
##
## $method
## [1] "misclass"
##
## attr("class")
## [1] "prune" "tree.sequence"
```

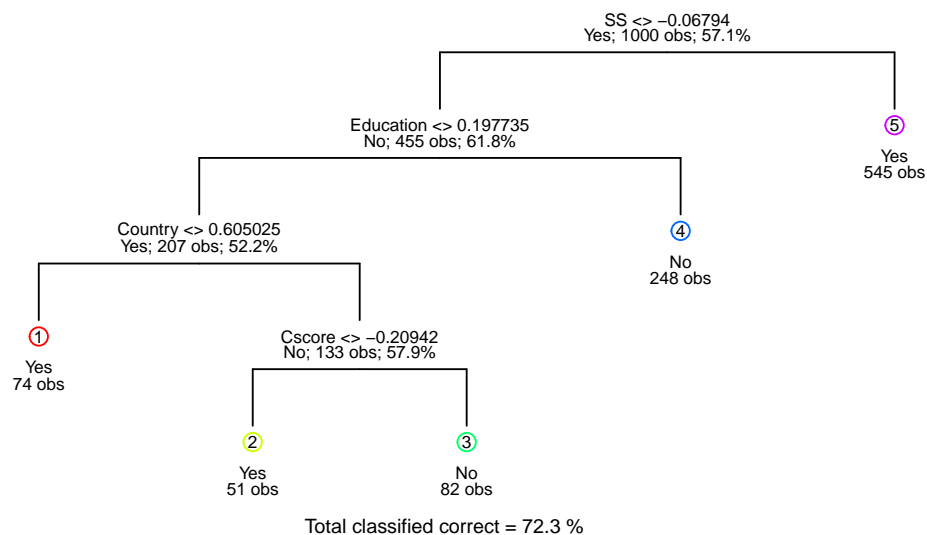
```
best.size = min(cv$size[cv$dev == min(cv$dev)])
best.size
```

```
## [1] 5
```

g.

```
pt.cv = prune.misclass(tree.drug, best=best.size)
draw.tree(pt.cv, nodeinfo=TRUE, cex = 0.55)
title("Pruned Classification Tree Built on Training Set")
```

## Pruned Classification Tree Built on Training Set



h. The True Positive Rate is 0.80 and the False Positive Rate is 0.47

```
tree.pred = predict(pt.cv, drug.test, type = 'class')

recent.test <- drug.test$recent_nicotine_use

error = table(recent.test, tree.pred)
error
```

```
##           tree.pred
## recent.test  No  Yes
##           No  208 188
##           Yes   96 393
```

```
TN = error[1]
FP = error[3]
FN = error[2]
TP = error[4]

#True Positive Rate = TP/(TP+FN)
paste('True Positive Rate: ', TP/(TP+FN))
```

```
## [1] "True Positive Rate: 0.803680981595092"
```

```
#False Positive Rate = FP/(FP+TN)
paste('False Positive Rate: ', FP/(FP+TN))
```

```
## [1] "False Positive Rate: 0.474747474747475"
```


- i. The variables that seem to be the most important are SS, Country, and Age.

```
set.seed(123)

boost.drug = gbm(ifelse(recent_nicotine_use == 'Yes', 1, 0)~., data=drug.train,
                  distribution = 'bernoulli', n.trees = 1000, shrinkage = 0.01)

summary(boost.drug)
```





```
##          var rel.inf
## SS          SS 26.214
## Country    Country 18.166
## Age        Age 15.498
## Cscore     Cscore 10.288
## Education  Education 7.365
## Oscore     Oscore 4.641
## Ascore     Ascore 3.504
## Gender     Gender 3.480
## Escore     Escore 3.270
## Impulsive  Impulsive 3.266
## Nscore     Nscore 2.814
## Ethnicity  Ethnicity 1.495
```

- j. The OOB estimate of error rate is 30%, There were 3 variables randomly considered at each split in the trees. 500 trees were used to fit the data. Despite some differences, the order of important variables are fairly similar between the boosting and random forest models.

```
set.seed(123)

rf.drug <- randomForest(recent_nicotine_use~., data=drug.train,
                        importance = TRUE)

rf.drug
```

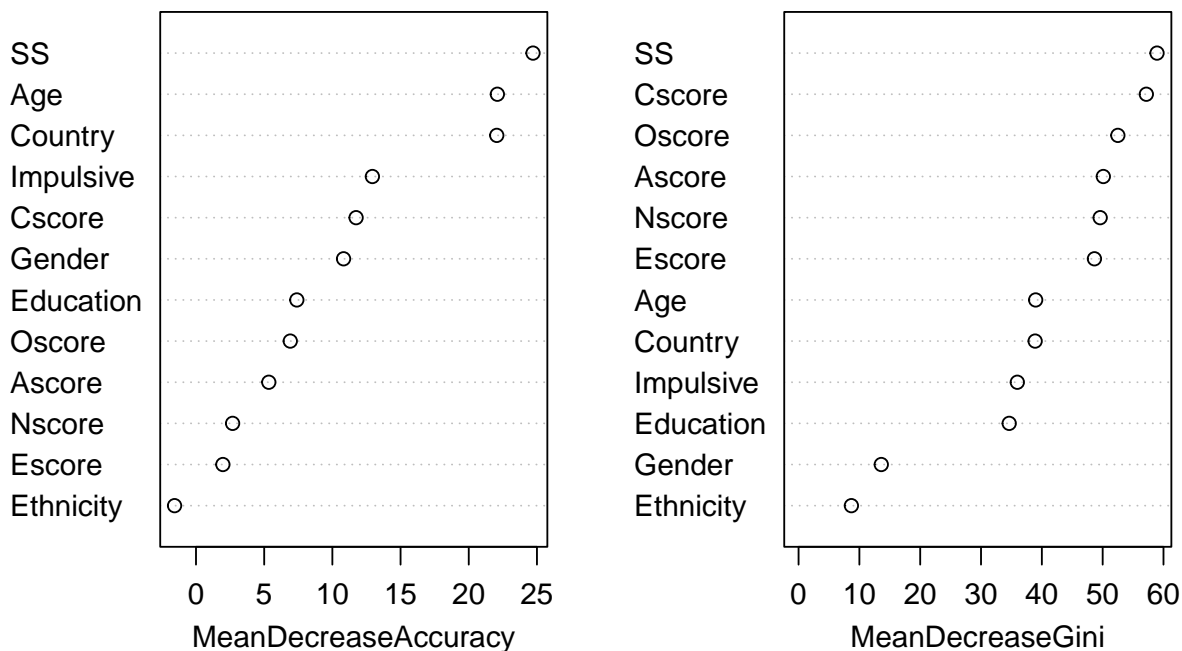
```
##
## Call:
## randomForest(formula = recent_nicotine_use ~ ., data = drug.train, importance = TRUE)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 3
##
##           OOB estimate of  error rate: 30%
## Confusion matrix:
##           No Yes class.error
## No  255 174      0.4056
## Yes 126 445      0.2207
```

```
importance(rf.drug)
```

	No	Yes	MeanDecreaseAccuracy	MeanDecreaseGini
## Age	15.4620	15.3910	22.097	38.984
## Gender	13.2275	2.2813	10.820	13.599
## Education	7.5800	2.9781	7.393	34.631
## Country	25.4077	5.0218	22.055	38.901
## Ethnicity	-1.3869	-0.8032	-1.572	8.683
## Nscore	-0.7845	4.2984	2.684	49.590
## Escore	2.3166	0.4504	1.969	48.645
## Oscore	4.5145	5.0452	6.913	52.499
## Ascore	5.8128	1.5615	5.339	50.107
## Cscore	11.4552	4.5539	11.739	57.178
## Impulsive	1.5397	13.7528	12.934	35.968
## SS	16.6944	17.4728	24.712	58.939

```
varImpPlot(rf.drug, sort=T, main='Variable Importance for rf.drug')
```

## Variable Importance for rf.drug



k. The fraction of people predicted to use nicotine recently who have in fact used nicotine recently is 0.5716.

```
prob.forest = predict(rf.drug, newdata=drug.test, type = 'prob')[, 'Yes']
pred.forest = ifelse(prob.forest >= 0.2, 'Yes', 'No')
```

```
rf.matrix = table(drug.test$recent_nicotine_use, pred.forest)
rf.matrix
```

```
##      pred.forest
##      No Yes
## No   43 353
## Yes  18 471
```

```
response.boost <- predict(boost.drug, newdata = drug.test, type = 'response')
```

```
## Using 1000 trees...
```

```
pred.boost = ifelse(response.boost >= 0.2, 'Yes', 'No')
```

```
boost.matrix = table(drug.test$recent_nicotine_use, pred.boost)
boost.matrix
```

```
##      pred.boost
##      No Yes
```

```
## No 24 372
## Yes 8 481
```

```
#forest - Fraction of predicted nicotine use who actually use nicotine.
rf.matrix[4] / (rf.matrix[3]+rf.matrix[4])
```

```
## [1] 0.5716
```